

Machine Learning for Computer Vision

TD2: Image segmentation

Département d'informatique

Ruiwen HE

Course Review: Image segmentation

1. What is Image Segmentation?

Image segmentation is the process of dividing an image into multiple meaningful regions. It plays a crucial role in computer vision by aiding in object recognition and image content understanding. It is widely applied in fields such as medical imaging, autonomous driving, and object detection.

2. Types of Image Segmentation Tasks:

- **Semantic Segmentation:** Classifies each pixel into a category without distinguishing between instances.
- **Instance Segmentation:** Differentiates between instances of the same class.
- **Panoptic Segmentation:** A combination of semantic and instance segmentation.

3. Traditional Image Segmentation Techniques:

- **Threshold-based Methods:** Separates foreground and background based on pixel intensity.
- **Edge-based Methods:** Detect object boundaries (e.g., Sobel, Canny edge detection).
- **Clustering-based Methods:** Segments by grouping similar pixels (e.g., K-Means, Gaussian Mixture Model).
- **Region-based Methods:** Groups neighboring pixels that share similar characteristics (e.g., Watershed Algorithm).

4. Deep Learning-based Image Segmentation:

- **Fully Convolutional Networks (FCNs):** Used for semantic segmentation.
- **U-Net:** Commonly used in medical imaging for precise localization.
- **Mask R-CNN:** Used for instance segmentation by detecting objects and generating pixel-level masks.
- **Panoptic FPN:** Combines semantic and instance segmentation into one framework.

5. Evaluation Metrics for Image Segmentation:

- **Pixel Accuracy:** Ratio of correctly classified pixels.
- **Intersection over Union (IoU):** Measures overlap between predicted and ground truth segmentation.
- **Dice Coefficient:** Similar to IoU, but assigns more weight to correctly classified pixels.

6. Applications of Image Segmentation:

- **Robotics:** Used in machine vision for object detection and navigation.
- **Medical Imaging:** Helps in disease diagnosis, organ segmentation, and surgical video annotation.
- **Autonomous Vehicles:** Critical for object detection and scene understanding.
- **Satellite Imaging:** Used to segment roads, buildings, and trees in aerial images.

Part I: Basics of Image Segmentation and Evaluation Metrics

Objective:

- Understand the fundamental concepts of image segmentation.
- Learn the types of segmentation tasks and evaluation metrics.

Exercise 1: Implement a function to calculate IoU and Dice Coefficient

- Intersection over Union (IoU): Measures the overlap between predicted segmentation and ground truth.

$$\text{IoU} = \frac{|A \cap B|}{|A \cup B|} = \frac{\sum_{i=1}^N y_i \hat{y}_i}{\sum_{i=1}^N (y_i + \hat{y}_i) - \sum_{i=1}^N y_i \hat{y}_i}$$

- y_i is the ground truth value for pixel i .
 - \hat{y}_i is the predicted value for pixel i .
 - N is the total number of pixels.
- Dice Coefficient: A measure of overlap between predicted and actual segmentation.

$$\text{Dice} = \frac{2|A \cap B|}{|A| + |B|} = \frac{2 \sum_{i=1}^N y_i \hat{y}_i}{\sum_{i=1}^N y_i + \sum_{i=1}^N \hat{y}_i}$$

- y_i is the ground truth value for pixel i .
 - \hat{y}_i is the predicted value for pixel i .
 - N is the total number of pixels.
- **Implementation:** Create two functions to calculate the Intersection over Union and the Dice Coefficient. Provide execution on examples of your choice. (hint: use libraries like numpy)
 - **Question:** How pixel accuracy can be misleading in imbalanced datasets ? (e.g., when the background dominates)
 - **Bonus:** Illustrate this with sample data.

Part II: Traditional Image Segmentation Techniques

Objective:

- Learn and implement traditional image segmentation techniques.
- Conduct a comparative analysis of these methods, highlighting their advantages and limitations.

Task:

- Implement Otsu's method for thresholding.
- Use Sobel and Canny edge detection for edge-based segmentation.
- Implement K-Means clustering for image segmentation.
- Use the Watershed algorithm for segmentation.

Exercise 2: Otsu Thresholding

- **Steps:**

1. **Read the Image**

Load the input image and convert it to grayscale if it is a color image:

$$I_{\text{gray}} = \text{cv2.cvtColor}(I_{\text{image}}, \text{cv2.COLOR_BGR2GRAY})$$

2. **Compute Histogram**

Calculate the histogram of the grayscale image to get the pixel intensity distribution:

$$\text{hist} = \text{cv2.calcHist}([I_{\text{gray}}], [0], \text{None}, [256], [0, 256])$$

3. **Otsu's Method**

Apply Otsu's method to determine the optimal threshold value T that minimizes the within-class variance:

$$T = \text{cv2.threshold}(I_{\text{gray}}, 0, 255, \text{cv2.THRESH_OTSU})[0]$$

4. **Thresholding**

Create a binary image by applying the threshold T :

$$I_{\text{binary}} = \begin{cases} 255, & \text{if } I_{\text{gray}} > T \\ 0, & \text{if } I_{\text{gray}} \leq T \end{cases}$$

5. **Display Results**

Show the original grayscale image and the binary image using a plotting library.

- **Implementation:** Implement the previous 5 steps. (hint: use matplotlib for visualization)
- **Question:** What are the strengths and limitations of Otsu's method ?

Exercise 3: Sobel Edge Detection

- **Steps:**

1. **Grayscale Conversion**

If the input image is a color image (RGB), convert it to grayscale, since edge detection is typically applied to single-channel images. The grayscale image I_{gray} can be computed as:

$$I_{\text{gray}} = 0.2989 I_R + 0.5870 I_G + 0.1140 I_B$$

where I_R , I_G , and I_B are the red, green, and blue channel intensities, respectively.

2. **Sobel Kernels**

Define the Sobel kernels for both horizontal and vertical gradients. These kernels are convolved with the grayscale image.

- **Horizontal Kernel (G_x):**

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

- **Vertical Kernel (G_y):**

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

3. **Convolution**

Perform convolution of the Sobel kernels with the grayscale image to compute the gradients in both horizontal (G_x) and vertical (G_y) directions:

$$G_x = I_{\text{gray}} * G_x \quad \text{and} \quad G_y = I_{\text{gray}} * G_y$$

where $*$ represents the convolution operation.

4. Gradient Magnitude

Compute the magnitude of the gradient at each pixel by combining the horizontal and vertical gradients:

$$G = \sqrt{G_x^2 + G_y^2}$$

5. Thresholding (optional)

Apply a threshold to the gradient magnitude to create a binary edge map. Pixels with values above the threshold T are considered edges:

$$G_{\text{thresholded}} = \begin{cases} 1, & \text{if } G > T \\ 0, & \text{if } G \leq T \end{cases}$$

- **Implementation:** Implement the previous 5 steps. (hint: use opencv Sobel and threshold functions, use matplotlib for visualization)
- **Question:** What are the strengths and limitations of Sobel's method ?

Exercise 4: Canny Edge Detection

■ Steps:

1. Grayscale Conversion

If the input image is a color image (RGB), convert it to grayscale, as Canny edge detection is performed on single-channel images. The grayscale image I_{gray} can be computed as:

$$I_{\text{gray}} = 0.2989 I_R + 0.5870 I_G + 0.1140 I_B$$

2. Gaussian Blurring

Apply a Gaussian filter to the grayscale image to reduce noise and avoid false edge detection:

$$I_{\text{blurred}} = G_{\sigma} * I_{\text{gray}}$$

where G_{σ} is the Gaussian kernel.

3. Gradient Calculation

Compute the gradients using Sobel operators to find the intensity and direction of edges:

$$G_x = I_{\text{blurred}} * G_x \quad \text{and} \quad G_y = I_{\text{blurred}} * G_y$$

4. Non-Maximum Suppression

Thin the edges by suppressing non-maximum values in the gradient direction. This step helps to maintain only the strongest edges.

5. Double Thresholding

Apply two thresholds (low and high) to determine strong and weak edges:

$$\text{Strong edges} \Rightarrow G > T_{\text{high}} \quad \text{Weak edges} \Rightarrow T_{\text{low}} < G \leq T_{\text{high}}$$

6. Edge Tracking by Hysteresis

Finalize the edges by suppressing weak edges that are not connected to strong edges.

- **Implementation:** Provide an implementation for previous steps. (hint: use opencv Canny function)
- **Questions:** What are the strengths and limitations of Canny Edge method ?

Exercise 5: K-Means Clustering for Image Segmentation

■ Steps:

1. Read and Reshape the Image

Load the input image and reshape it into a two-dimensional array where each pixel is represented by its RGB values:

$$\text{pixels} = I_{\text{image}}.\text{reshape}(-1, 3)$$

2. Apply K-Means Clustering

Use the K-Means algorithm to cluster the pixel values into k clusters:

```
kmeans = KMeans(n_clusters = k, random_state = 42)
```

Fit the K-Means model to the pixel data:

```
kmeans.fit(pixels)
```

3. Replace Each Pixel with its Cluster Centroid

Replace each pixel in the image with the centroid of its assigned cluster:

```
segmented_image = kmeans.cluster_centers[kmeans.labels_].reshape(I_image.shape).astype(np.uint8)
```

4. Display Results

Show the original image and the segmented image using a plotting library.

- **Implementation:** Provide an implementation for previous steps.
- **Questions:** What are the advantages and disadvantages of using K-Means for image segmentation compared to other methods? Name at least one other clustering method and compare with K-Means.

Exercise 6: Watershed Algorithm for Image Segmentation

■ Steps:

1. Read and Preprocess the Image

Load the image and convert it to grayscale if it is in color. Apply Gaussian smoothing to reduce noise:

```
I_gray = cv2.cvtColor(I, cv2.COLOR_BGR2GRAY)
```

```
I_smoothed = cv2.GaussianBlur(I_gray, (5, 5), 0)
```

2. Compute the Gradient

Compute the gradient of the image to identify the edges:

```
I_gradient = cv2.Laplacian(I_smoothed, cv2.CV_64F)
```

3. Thresholding and Dilation

Apply a binary threshold to the gradient image and perform morphological dilation to enhance the markers:

```
ret, markers = cv2.threshold(I_gradient, T, 255, cv2.THRESH_BINARY_INV)
```

```
markers = cv2.dilate(markers, None, iterations = 3)
```

4. Watershed Algorithm

Apply the watershed algorithm to segment the image:

```
I_markers = cv2.watershed(I, markers)
```

Mark the boundaries of the segmented regions:

```
I[I_markers == -1] = [0, 0, 255] (Red for boundaries)
```

5. Display Results

Show the segmented image with the boundaries:

```
plt.imshow(I)
```

- **Implementation:** Provide an implementation for previous steps. (hint: use opencv functions GaussianBlur, Laplacian, threshold, dilate and watershed)
- **Question:** What are the advantages and limitations of the watershed algorithm in image segmentation?

Part III: Deep Learning-Based Image Segmentation Methods

Objective:

- Implement various deep learning models for semantic segmentation, instance segmentation, and panoptic segmentation.

Exercise 7: Image Segmentation using Pre-trained FCN Model

■ Steps:

1. Import Libraries

Import the necessary libraries for image processing and deep learning:

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from torchvision import models
5 import torch
6 from torchvision import transforms
```

2. Load the Pre-trained FCN Model

Load a pre-trained FCN model from PyTorch's torchvision library:

```
1 # Load pre-trained FCN model
2 model = models.segmentation.fcn_resnet101(pretrained=True)
3 model.eval() # Set the model to evaluation mode
```

3. Preprocess the Input Image

Load and preprocess the input image, resizing it and normalizing:

```
1 # Load and preprocess the image
2 image = cv2.imread('image.jpg')
3 image_resized = cv2.resize(image, (640, 480)) # Resize to match the
         input size
4 input_tensor = transforms.ToTensor()(image_resized).unsqueeze(0) #
         Add batch dimension
```

4. Perform Segmentation

Use the model to predict segmentation:

```
1 # Perform segmentation
2 with torch.no_grad():
3     output = model(input_tensor)['out'][0] # Get the output from
         the model
4     output_predictions = output.argmax(0) # Get the predicted class
         for each pixel
```

5. Visualize the Segmentation Result

Display the original image and the segmentation result:

```
1 # Visualize the results
2 plt.figure(figsize=(10, 5))
3 plt.subplot(1, 2, 1)
4 plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
5 plt.title('Original Image')
6 plt.axis('off')
7
8 plt.subplot(1, 2, 2)
9 plt.imshow(output_predictions.cpu().numpy(), cmap='jet', alpha=0.5)
10 plt.title('Segmentation Result')
11 plt.axis('off')
12
13 plt.show()
```

- **Question:** What are the advantages of using pre-trained models for image segmentation tasks?

Exercise 8: U-Net for Image Segmentation

- **Overview:** U-Net is a convolutional neural network designed for biomedical image segmentation. It employs a symmetric architecture that enables precise localization while maintaining context through skip connections.
- **Architecture:**
 - U-Net consists of a contracting path (encoder) to capture context and a symmetric expanding path (decoder) for precise localization.
 - The network uses skip connections to combine features from the encoder with the decoder, enhancing the model's ability to segment images accurately.
- **Code Example:**

```
1 import torch
2 import torchvision.transforms as transforms
3 from torchvision import models
4 from PIL import Image
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 # Load the pre-trained U-Net model
9 model = models.segmentation.unet(pretrained=True)
10 model.eval()
11
12 # Load and preprocess the image
13 image = Image.open('image.jpg').convert('RGB')
14 transform = transforms.Compose([
15     transforms.Resize((256, 256)),
16     transforms.ToTensor(),
17     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
18         0.225]),
19 ])
20 input_tensor = transform(image).unsqueeze(0) # Add batch dimension
21
22 # Perform inference
23 with torch.no_grad():
24     output = model(input_tensor)['out'][0] # Get the output
25     output_predictions = output.argmax(0) # Get the class with the
26         highest score
27
28 # Display the results
29 plt.subplot(1, 2, 1)
30 plt.title("Original Image")
31 plt.imshow(image)
32
33 plt.subplot(1, 2, 2)
34 plt.title("U-Net Segmentation")
35 plt.imshow(output_predictions.cpu().numpy(), cmap='jet')
36 plt.colorbar()
37 plt.show()
```

- **Code Explanation:**
 - **U-Net Model:** This code example demonstrates how to load a pre-trained U-Net model. Note that PyTorch's 'torchvision' library may not directly provide U-Net, so you might need to use alternative implementations like 'segmentation_models.pytorch'.
 - **Image Preprocessing:** The input image is resized to 256×256 pixels and normalized.
 - **Inference:** The model is used to perform inference, and the predicted output is obtained.
 - **Results Display:** The original image and its corresponding segmentation result are displayed.
- **Question:** What are the advantages of using U-Net over traditional segmentation methods, and how do skip connections enhance performance?

Exercise 9: Mask R-CNN for Instance Segmentation

- **Overview:** Mask R-CNN is an extension of Faster R-CNN that adds a branch for predicting segmentation masks on each Region of Interest (RoI), in parallel with the existing branch for classification and bounding box regression.
- **Architecture:**
 - Mask R-CNN consists of a backbone network (such as ResNet) for feature extraction, a Region Proposal Network (RPN) for proposing candidate object bounding boxes, and a branch for predicting segmentation masks.
 - The network uses a Fully Convolutional Network (FCN) for mask prediction, allowing it to output high-resolution segmentation masks.
- **Code Example:**

```
1 import torch
2 from torchvision import models
3 from PIL import Image
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import torchvision.transforms as T
7
8 # Load the pre-trained Mask R-CNN model
9 model = models.detection.maskrcnn_resnet50_fpn(pretrained=True)
10 model.eval()
11
12 # Load and preprocess the image
13 image = Image.open('image.jpg').convert('RGB')
14 transform = T.Compose([
15     T.ToTensor(),
16 ])
17 input_tensor = transform(image) # Convert the image to a tensor
18
19 # Perform inference
20 with torch.no_grad():
21     predictions = model([input_tensor]) # Get predictions
22
23 # Visualize the results
24 plt.imshow(image)
25 ax = plt.gca()
26
27 # Draw bounding boxes and masks
28 for i in range(len(predictions[0]['boxes'])):
29     box = predictions[0]['boxes'][i].numpy()
30     score = predictions[0]['scores'][i].item()
31     mask = predictions[0]['masks'][i, 0].mul(255).byte().numpy()
32
33     if score > 0.5: # Confidence threshold
34         ax.add_patch(plt.Rectangle((box[0], box[1]), box[2]-box[0], box
35                                     [3]-box[1],
36                                     fill=False, edgecolor='red',
37                                     linewidth=2))
38         plt.imshow(mask, alpha=0.5) # Overlay mask
39
40 plt.title("Mask R-CNN Segmentation")
41 plt.axis('off')
42 plt.show()
```

- **Code Explanation:**
 - **Mask R-CNN Model:** This code snippet demonstrates how to load a pre-trained Mask R-CNN model using PyTorch's 'torchvision' library.

- **Image Preprocessing:** The input image is transformed into a tensor suitable for the model.
 - **Inference:** The model performs inference on the input image, outputting bounding boxes, scores, and masks for detected objects.
 - **Results Visualization:** The original image is displayed with overlaid bounding boxes and masks for detected objects, filtered by a confidence threshold.
- **Question:** What are the benefits of using Mask R-CNN for instance segmentation compared to traditional methods, and how does it manage the trade-off between speed and accuracy?

Exercise 10: Panoptic FPN Segmentation

- **Overview:** The Panoptic Feature Pyramid Network (Panoptic FPN) is designed to achieve high-quality object detection and segmentation by leveraging multi-scale feature representations. This approach performs both instance segmentation and semantic segmentation in a unified framework.
- **Steps:**
 1. **Load Pre-trained Model**
Import the necessary libraries and load a pre-trained Panoptic FPN model, typically trained on the COCO dataset.
 2. **Prepare Input Image**
Preprocess the input image to match the expected dimensions and normalize pixel values. The input should be a tensor with shape $[1, 3, H, W]$.
 3. **Perform Inference**
Pass the input image through the model to obtain predictions, including bounding boxes, segmentation masks, and class labels.
 4. **Post-Processing**
Apply post-processing steps to refine the predictions, such as thresholding for confidence scores and applying non-maximum suppression to eliminate duplicate boxes.
 5. **Visualize Results**
Visualize the segmentation results by overlaying the predicted masks and bounding boxes on the original image.
- **Code Example:**

```

1 import torch
2 from torchvision import models
3 import torchvision.transforms as T
4 from PIL import Image
5 import matplotlib.pyplot as plt
6
7 # Load the pre-trained Panoptic FPN model
8 model = models.detection.maskrcnn_resnet50_fpn(pretrained=True)
9 model.eval()
10
11 # Load and preprocess the input image
12 input_image = Image.open('path/to/your/image.jpg').convert("RGB")
13 transform = T.Compose([
14     T.Resize((800, 800)),
15     T.ToTensor(),
16 ])
17 input_tensor = transform(input_image).unsqueeze(0) # Add batch
18           dimension
19
20 # Perform inference
21 with torch.no_grad():
22     predictions = model(input_tensor)
23
24 # Visualize the results
25 plt.figure(figsize=(12, 8))
26 plt.imshow(input_image)

```

```

26
27 # Overlay masks and boxes
28 for i in range(len(predictions[0]['boxes'])):
29     box = predictions[0]['boxes'][i].cpu().numpy()
30     score = predictions[0]['scores'][i].cpu().numpy()
31     mask = predictions[0]['masks'][i, 0].cpu().numpy()
32
33     if score > 0.5: # Confidence threshold
34         plt.imshow(mask, alpha=0.5) # Overlay mask
35         plt.gca().add_patch(plt.Rectangle(
36             (box[0], box[1]), box[2]-box[0], box[3]-box[1],
37             fill=False, color='red', linewidth=2
38         ))
39
40 plt.axis('off')
41 plt.title("Panoptic FPN Segmentation Results")
42 plt.show()

```

■ Code Explanation:

- The code imports necessary libraries, including PyTorch and torchvision, for loading and using the pre-trained model.
- It loads the pre-trained Panoptic FPN model using 'maskrcnn_resnet50_fpn', which combines object detection and instance segmentation.
- An input image is loaded and preprocessed, resized to 800×800 pixels and converted to a tensor.
- The model performs inference on the input tensor, providing predictions that include bounding boxes, segmentation masks, and confidence scores.
- The results are visualized by overlaying the predicted masks and bounding boxes onto the original image, applying a confidence threshold to filter out less certain predictions.

- **Question:** What are the key differences between FPN and the two previous methods ?

Part IV: Training on Your Own Dataset (BONUS)

Objective:

Learn how to prepare a custom dataset for image segmentation and train a deep learning model (e.g., U-Net or Mask R-CNN) on this dataset.

Content Overview:

■ Dataset Preparation:

- How to format your dataset for training.
- Creating images and corresponding masks.
- Splitting the dataset into training and validation sets.

■ Model Setup:

- Choosing a model architecture (e.g., U-Net, Mask R-CNN).
- Configuring hyperparameters.

■ Training the Model:

- Setting up the training loop.
- Monitoring training and validation metrics.

■ Evaluating Results:

- Evaluating the trained model on the test set.
- Visualizing the predictions.